



# RUST CHINA CONF 2020

首届中国 Rust 开发者大会

2020.12.26-27 深圳

# 用 Rust 设计高性能 JIT 执行引擎

周鹤洋 @ Rust China Conf 2020

# 关于我

- GitHub @losfair
- 南京航空航天大学 2018 级本科生
- 编译 / OS / VM / 微架构
- 从 2017 年开始使用 Rust

# Projects that I worked on

- Wasmer / 跨平台 WebAssembly Runtime
- kernel-wasm / Linux 内核态 WebAssembly SFI 运行环境
- Violet / 函数式设计的超标量 RISC-V CPU
- VMesh / 去中心化 Mesh 路由协议
- Wavelet / DAG 分布式账本
- FlatMk / Capability-based microkernel in Rust
- ... and a lot of toy projects

Let's begin

# Just-In-Time compilation

- “即时编译”
- 运行时生成代码并执行

# Where are JITs used?

- 虚拟机
- 动态链接器
- Linux 内核

# Virtual machines

- JVM / CLR / V8 / Wasmer / Wasmtime / CPython etc.

Why?

目标语言规范与硬件规范不一致时，需要使用虚拟机 (VM) 运行

例如：

- 动态特性
- 硬件层面难以实现的沙盒特性 ( SFI )
- 不同的指令集 (二进制翻译)



# VM 技术

- 简单解释器
  - wasmi
- 优化的解释器
  - CPython, BEAM, wasm3
  - Threaded interpreter, fused operations, inline caching
- 即时编译 / Just-In-Time (JIT) Compilation
  - JVM, CLR, V8, LuaJIT, Wasmer, Wasmtime

# Why do we need JITs?

- 动态优化代码
- 高效支持语言的动态特性 / 安全性要求
- 初始化 (动态链接器 & Static call 机制)

# Dynamic optimization

# JIT 优化级别

JavaScriptCore	LLInt	Baseline JIT	DFG JIT	FTL JIT
V8	Ignition	-	-	TurboFan
Wasmer	-	Singlepass	Cranelift	LLVM


解释执行

JIT 级别 1

JIT 级别 2

JIT 级别 3

Profile & compile




解释执行

JIT 级别 1


JIT 级别 2

JIT 级别 3

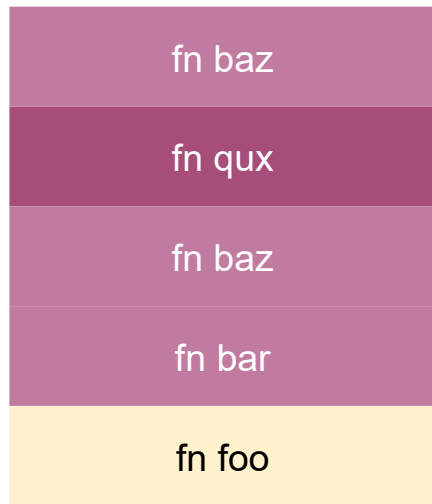


Deoptimize

# 切换优化级别：栈上替换 / On-Stack Replacement



当前调用栈



新的调用栈

解释执行

JIT 级别 1

JIT 级别 2

JIT 级别 3

# OSR in Wasmer

Su Engine @ <https://github.com/wasmerio/wasmer/pull/489>

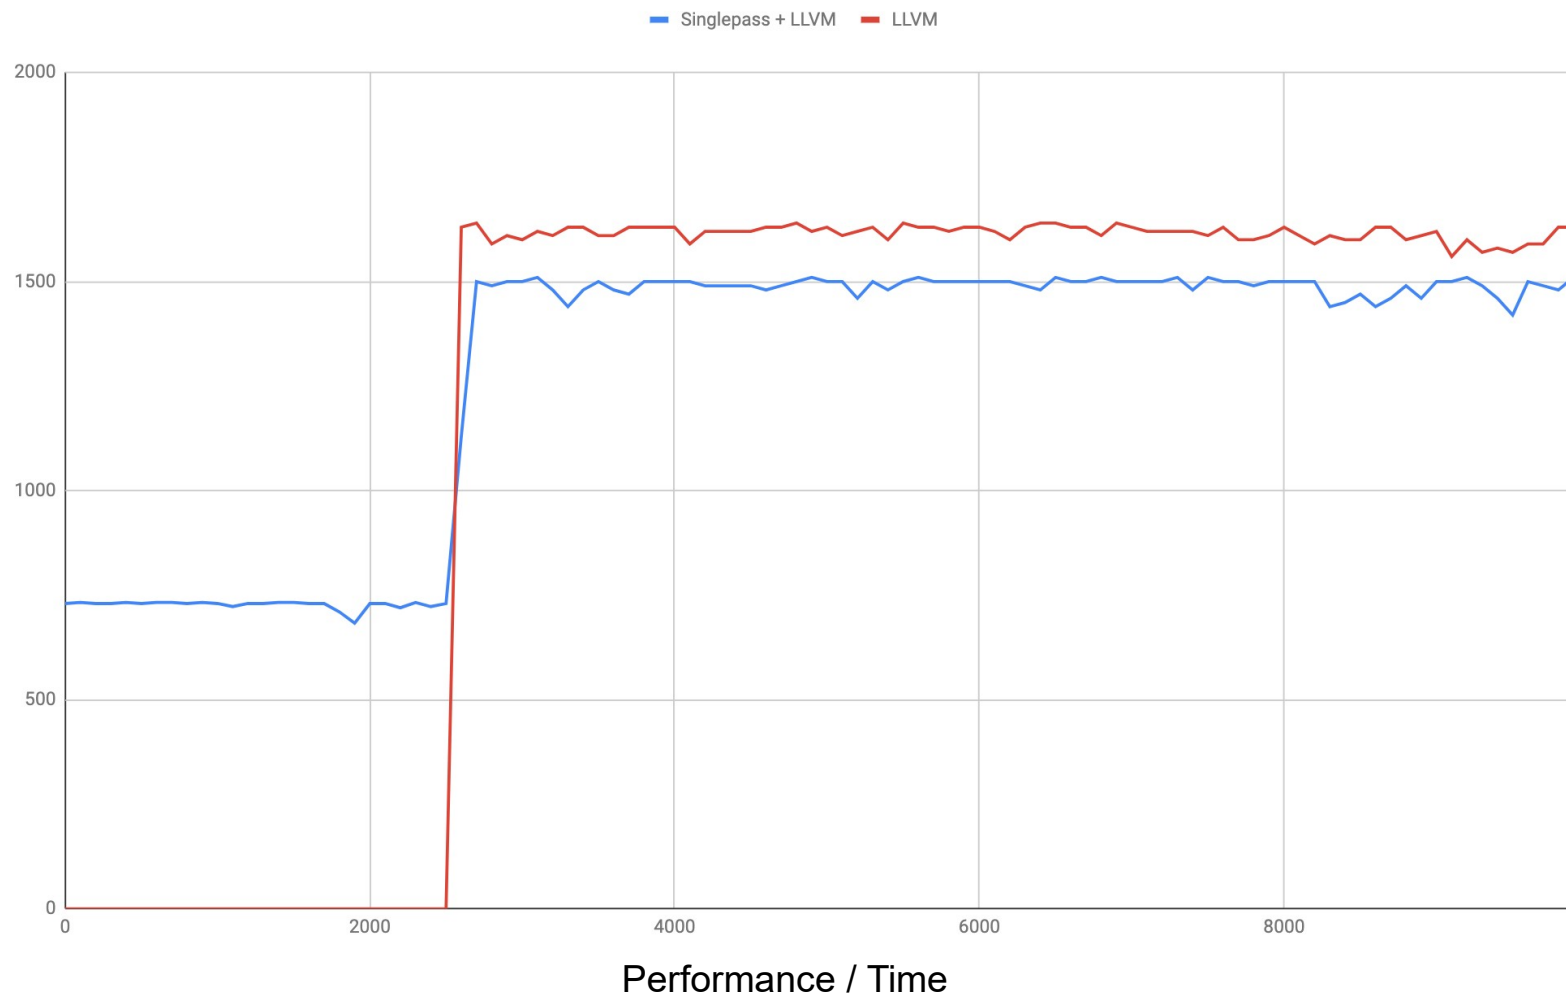
My work!

```
pub unsafe fn read_stack(  
    msm: &ModuleStateMap,  
    code_base: usize,  
    mut stack: *const u64,  
    initially_known_registers: [Option<u64>; 24],  
    mut initial_address: Option<u64>,  
) -> ExecutionStateImage {
```

OSR exit

```
pub unsafe fn invoke_call_return_on_stack(  
    msm: &ModuleStateMap,  
    code_base: usize,  
    image: &InstanceImage,  
    vmctx: &mut Ctx,  
) -> u64 {
```

OSR entry





# 内联缓存

## JavaScript

```
for(let x of list) {  
    document.write(x); // Method lookup  
}
```

## RISC-V 二进制翻译

11040:	f9843603	ld	a2,-104(s0) # MMU lookup
11044:	00078593	mv	a1,a5
11048:	00070513	mv	a0,a4
1104c:	000680e7	jalr	a3 # MMU lookup & translation
lookup			
11050:	00050793	mv	a5,a0
11054:	02f04463	bgtz	a5,1107c
<core_list_mergesort+0x168>			
11058:	fe843783	ld	a5,-24(s0) # MMU lookup

## rvjit-aa64: 从内存加载数据（内联缓存快速路径）

```
// Compute effective address
ld_simm16(&mut self.a, trash_reg_w() as _, imm);
dynasm!(self.a
    ; .arch aarch64
    ; add X(trash_reg_w() as u32), X(trash_reg_w() as u32), X(rs as u32)
);

let lower_bound_slot = self.alloc_rtslot(std::u64::MAX);
let upper_bound_slot = self.alloc_rtslot(0);
let reloff_slot = self.alloc_rtslot(std::u64::MAX);

self.load_rtslot_pair(lower_bound_slot, 30, temp1_reg());

// Check bounds
dynasm!(self.a
    ; .arch aarch64

    // Lower bound
    ; cmp X(trash_reg_w() as u32), x30
    ; b.lo >fallback

    // Upper bound
    ; cmp X(trash_reg_w() as u32), X(temp1_reg())
    ; b.hs >fallback
);
```

## rvjit-aa64: 从内存加载数据（慢速路径）

```
let addr = self.read_register(pp.rs as _) + (pp.rs_offset as i64 as u64);
debug!("load/store miss. target = 0x{:016x}", addr);

let (target_base_v, target_section) = match registry.lookup_section(addr) {
    Ok(x) => x,
    Err(_) => return Err(ExecError::BadLoadStoreAddress),
};

if pp.is_store && !target_section.get_flags().contains(SectionFlags::W) {
    return Err(ExecError::BadLoadStoreFlags);
}

if !pp.is_store && !target_section.get_flags().contains(SectionFlags::R) {
    return Err(ExecError::BadLoadStoreFlags);
}

let target_base_real = target_section.get().as_ptr() as u64;
let target_len = target_section.get().len() as u64;
```

Memory safety


# JIT 内存机制

- 空指针检查
- 访问越界检查


简单的实现方法： `compare-and-branch`

如何减少开销？

```
cmp $0, %rdi
je null_pointer_exception
mov %rdi, 16(%rsp)
...
null_pointer_exception:
    call host_npe_handler
    ...
```



```
mov %rdi, 16(%rsp)
...
Trap!
sigsegv_handler:
...
call host_npe_handler
...
```

A horizontal arrow points from the instruction `mov %rdi, 16(%rsp)` to the label `sigsegv_handler:`. Above the arrow, the word `Trap!` is written in a reddish-purple color.

# 采用 Unix 信号处理同步异常

```
// Allow handling OOB with signals on all architectures
register(&mut PREV_SIGSEGV, libc::SIGSEGV);

// Handle `unreachable` instructions which execute `ud2` right now
register(&mut PREV_SIGILL, libc::SIGILL);

// x86 uses SIGFPE to report division by zero
if cfg!(target_arch = "x86") || cfg!(target_arch = "x86_64") {
    register(&mut PREV_SIGFPE, libc::SIGFPE);
}

// On ARM, handle Unaligned Accesses.
// On Darwin, guard page accesses are raised as SIGBUS.
if cfg!(target_arch = "arm") || cfg!(target_os = "macos") {
    register(&mut PREV_SIGBUS, libc::SIGBUS);
}
```



JIT in Linux kernel

# JIT in Linux kernel: eBPF

- 允许用户代码安全介入内核的机制
- Interpreter + JIT

# JIT in Linux kernel: Static call (Linux 5.10)

- 动态生成直接跳转指令序列替代间接跳转
- 比 Retpoline 更高效的 Spectre v2 缓解机制

```
static void __ref __static_call_transform(void *insn, enum insn_type type, void *func)
{
    int size = CALL_INSN_SIZE;
    const void *code;

    switch (type) {
    case CALL:
        code = text_gen_insn(CALL_INSN_OPCODE, insn, func);
        break;

    case NOP:
        code = ideal_nops[NOP_ATOMIC5];
        break;

    case JMP:
        code = text_gen_insn(JMP32_INSN_OPCODE, insn, func);
        break;

    case RET:
        code = text_gen_insn(RET_INSN_OPCODE, insn, func);
        size = RET_INSN_SIZE;
        break;
    }

    if (memcmp(insn, code, size) == 0)
        return;

    if (unlikely(system_state == SYSTEM_BOOTING))
        return text_poke_early(insn, code, size);

    text_poke_bp(insn, code, size, NULL);
}
```

# JIT in Linux kernel: Static call (Linux 5.10)

```
some_func:
    jmp *%rax
    ...
```

```
some_func:
    call L2

L1:
    lfence

    jmp L1

L2:
    mov %rax, (%rsp)
    ret
    ...
```

```
some_func:
    call _trampoline
    ...

_trampoline: # Rewritten at runtime
    .byte 0xe9 # Relative JMP
    .dword (target - _trampoline)
```

Original


Retpoline

Static call

Should I use a JIT?

# 执行效率 & 工程复杂性

- 执行效率: JIT >>> 优化的解释器 >> 简单解释器
- 工程复杂性: JIT >>> 优化的解释器 > 简单解释器



# Complexity -> Bugs

LPE bugs discovered in Linux eBPF JIT:

- CVE-2020-8835
- CVE-2020-27194

# Implementing JIT in Rust



Thanks!



# RUST CHINA CONF 2020

首届中国 Rust 开发者大会

2020.12.26-27 深圳