



The Design of A High-performance Tracing Library in Rust

Zhenchi Zhong (钟镇炽) @  PingCAP




Who am I

- Infrastructure Engineer in PingCAP
 - PingCAP
 - TiDB
 - Distributed transactional relational database
 - HTAP, MySQL Compatibility...
 - TiKV
 - Distributed transactional key-value database
 - Support TiDB as storage engine
 - Developed in **Rust**
 - Graduated from Cloud Native Computing Foundation (CNCF)
- Mainly work on **tracing** for TiDB & TiKV



About Tracing

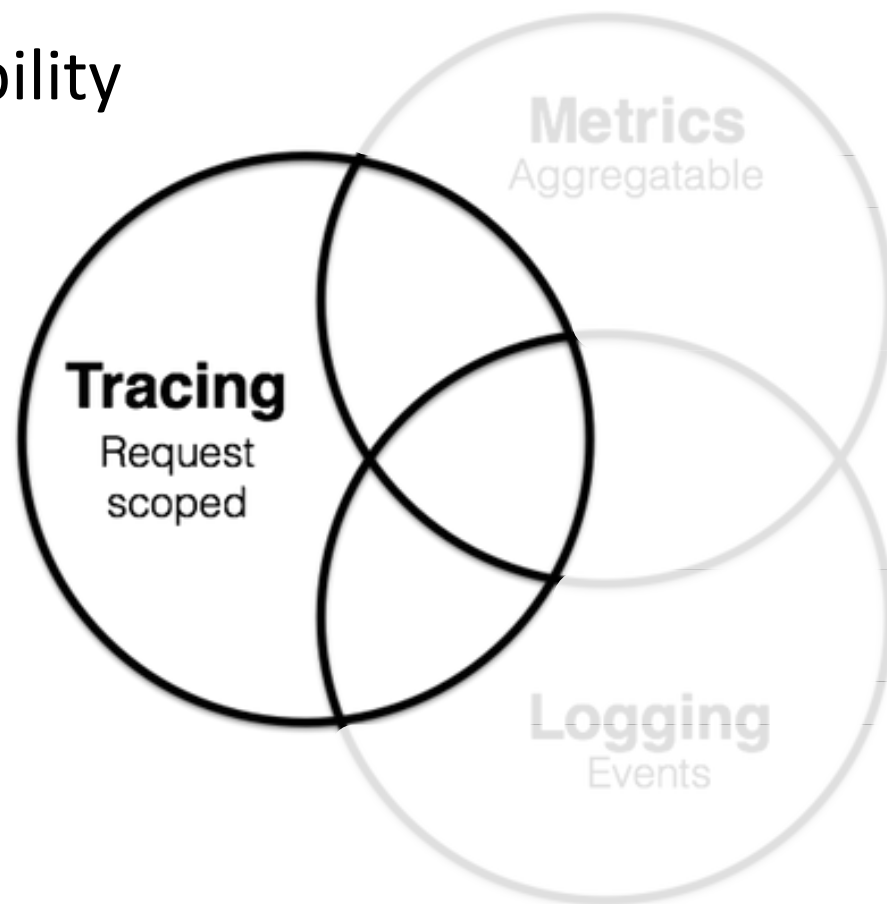
- One of the three pillars of observability





About Tracing

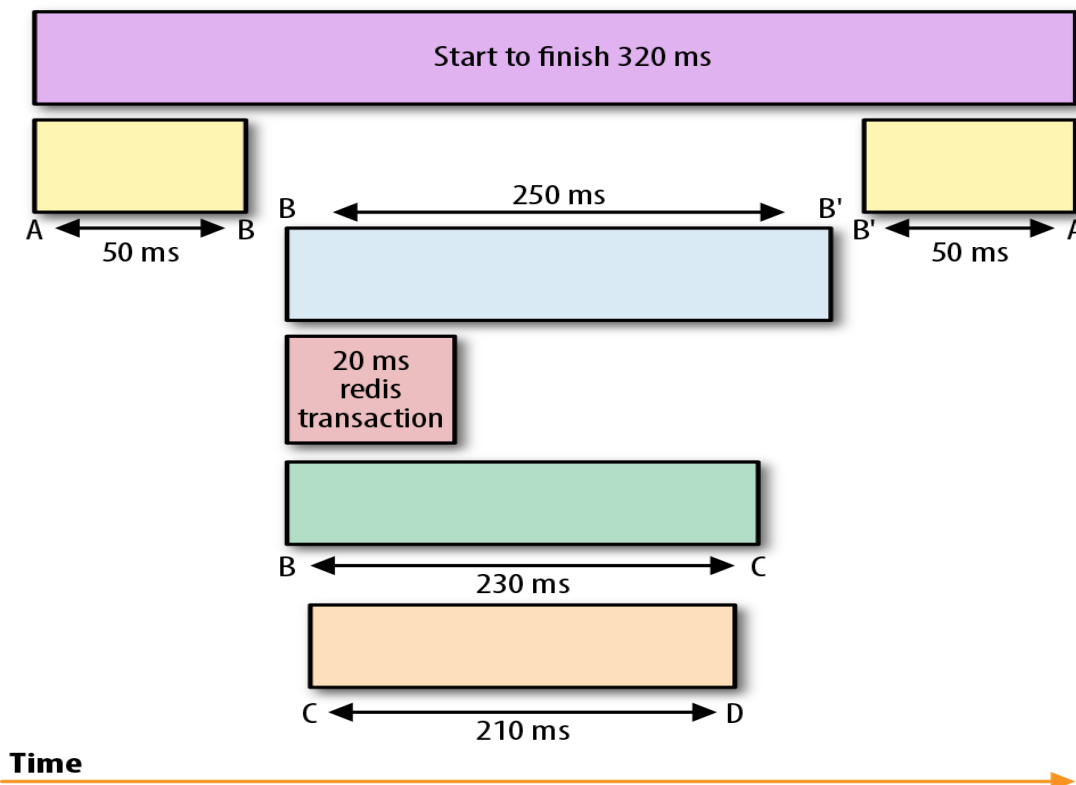
- One of the three pillars of observability
- Request scoped
- Use cases
 - Profiling
 - Debugging
 - Latency jitter analysis





About Tracing

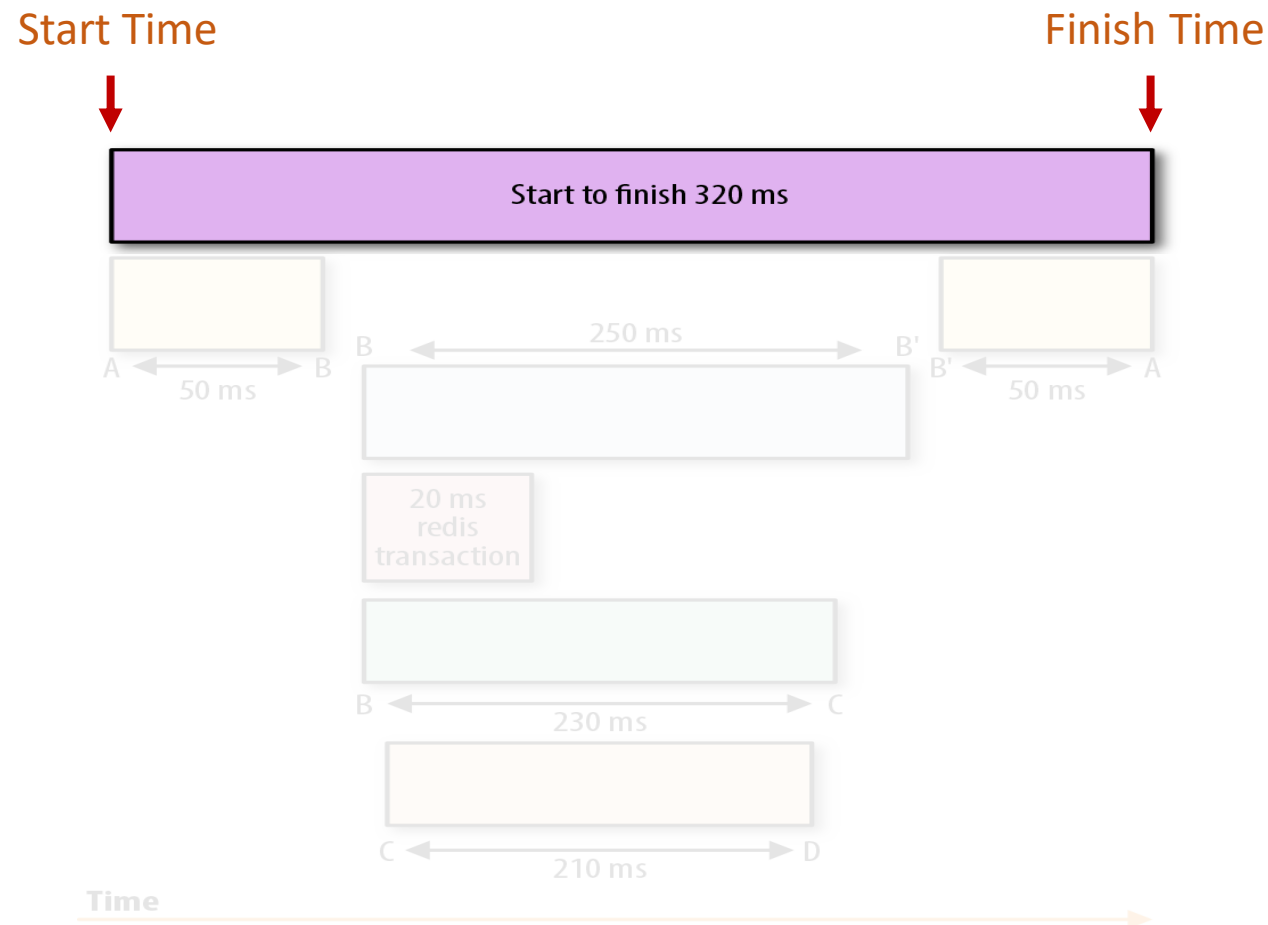
- Trace requests: HTTP, RPC...
- Across components
 - Redis
 - RDBMS
 - ...





About Tracing


- Span
 - Start time
 - Finish time
 - Operation name





About Tracing

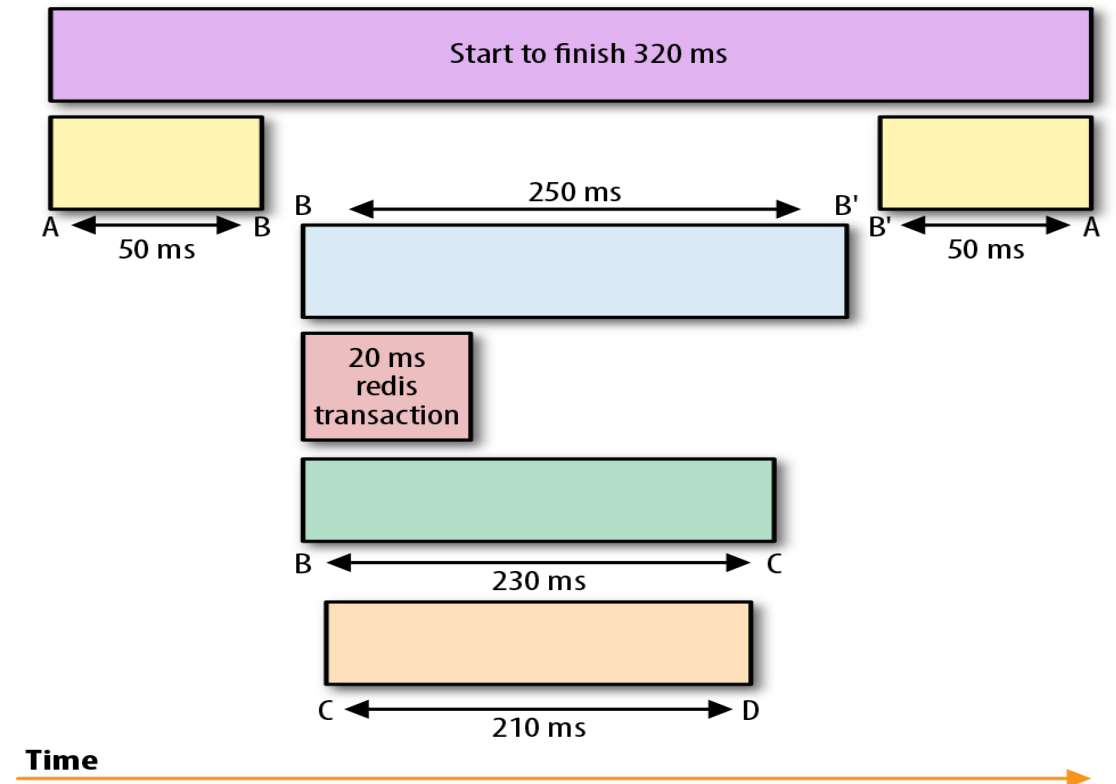
- Reference
 - Dependency
 - Parent-Child





About Tracing

- Reference
 - Dependency
 - Parent-Child







Early Try

- Plan
 - Choose a library from open-source community
 - Tokio tracing: <https://github.com/tokio-rs/tracing.git>
 - Rustracing: <https://github.com/sile/rustracing.git>
 - Apply it
 - Happy ending





Early Try

- Plan
 - Choose a library from open-source community
 - Tokio tracing: <https://github.com/tokio-rs/tracing.git>
 - Rustracing: <https://github.com/sile/rustracing.git>
 - Apply it
 - Happy ending 
- Reality
 - QPS is reduced by 50% if enabled 



Expected Library

- High performance (Minimal overhead)
- Compatibility with OpenTracing
 - Jaeger, Datadog...




Expected Library


- High performance (Minimal overhead)
- Compatibility with OpenTracing
 - Jaeger, Datadog...
- Minitrace
 - <https://github.com/tikv/minitrace-rust.git>



20ns!





Per Span



Micro benchmark (lower is better)



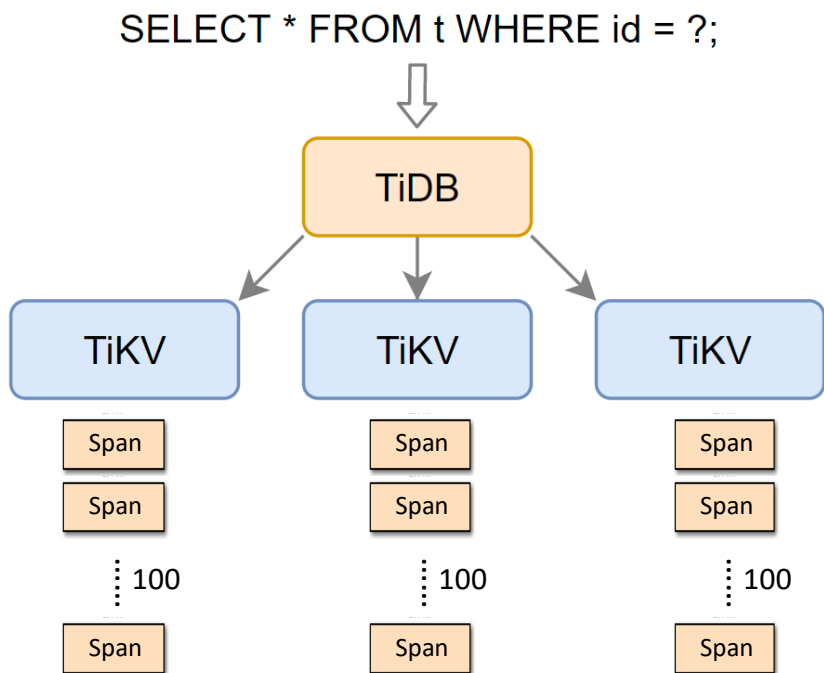
Overhead < 5%!



Integration benchmark (higher is better)



Overhead < 5%!



QPS

200 K

150 K

100 K

50 K

0 K

No tracing

Minitrace

Rustracing

> 50%

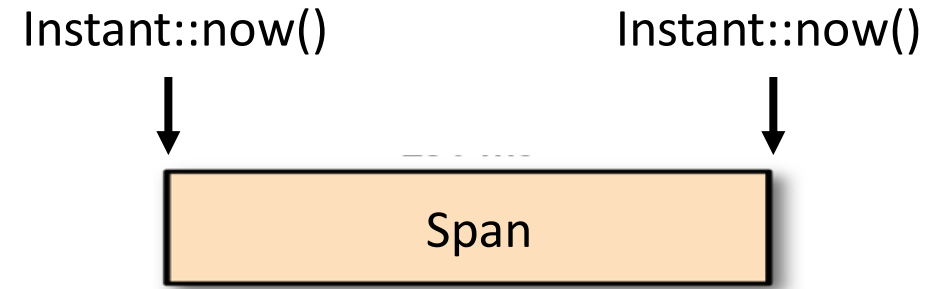


Integration benchmark (higher is better)



Naïve Implementation


- Timing approach
 - `std::time::SystemTime::now()`?
 - `std::time::Instant::now()`
 - One at started, one at finishing





Naïve Implementation


- Timing approach
 - `std::time::Instant::now()`
 - One at started, one at finishing
- A shared collector
 - From multiple threads





Naïve Implementation

- Timing approach
 - `std::time::Instant::now()`
 - One at started, one at finishing
- A shared collector
 - From multiple threads
 - `std::sync::mpsc::channel`

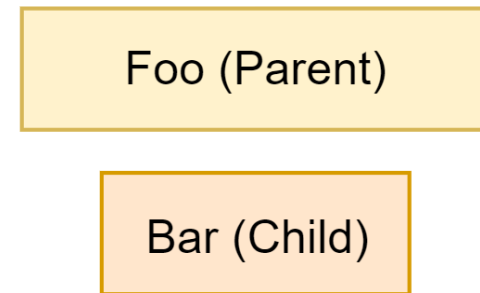




Naïve Implementation

- Timing approach
 - `std::time::Instant::now()`
 - One at started, one at finishing
- A shared collector
 - From multiple threads
 - `std::sync::mpsc::channel`
- An explicit context
 - Build references

```
fn foo(                ) {  
    bar( );  
}  
  
fn bar(                ) {  
  
}
```



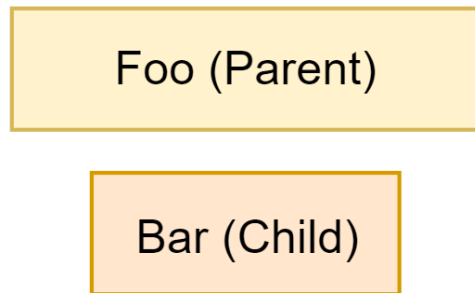


Naïve Implementation

- Timing approach
 - `std::time::Instant::now()`
 - One at started, one at finishing
- A shared collector
 - From multiple threads
 - `std::sync::mpsc::channel`
- An explicit context
 - Build references
 - Update context when creating spans
 - Retrieve parent span from context

```
fn foo(ctx: &mut Context) {  
    let span = ctx.create_span();  
    bar(ctx);  
}
```

```
fn bar(ctx: &mut Context) {  
    let span = ctx.create_span();  
}
```





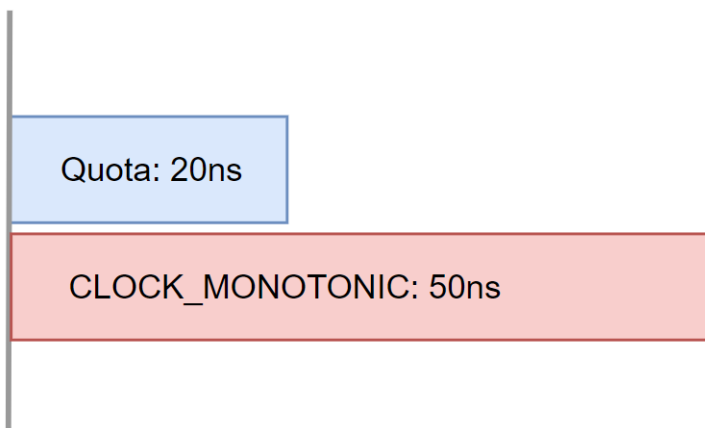

How to Improve

- Expected impact $< 5\%$
 - $< 20\text{ns}/\text{span}$



How to Improve: Timing



- Overhead of Timing
 - `std::time::Instant::now()`
 - Identical to `clock_gettime(CLOCK_MONOTONIC, ...)` in Linux
 - $25\text{ ns} \times 2$





How to Improve: Timing

- An Option of Low Precision
 - `coarsetime::Instant::now()`
 - Identical to `clock_gettime(CLOCK_MONOTONIC_COARSE, ...)` in Linux
 - $5\text{ ns} \times 2$






How to Improve: Timing


- An Option of Low Precision

- `coarsetime::Instant::now()`
- Identical to `clock_gettime(CLOCK_MONOTONIC_COARSE, ...)` in Linux
- $5\text{ ns} \times 2$
- Kernel jiffy precision (4ms by default)






How to Improve: Timing



High Precision





Precision of 4ms



How to Improve: Timing

- Better Choice on **x86/x64**
 - TimeStampCounter register
 - Increase per tick
 - **8 ns × 2**
 - Nanoseconds precision




```
fn tsc() -> u64 {  
    #[cfg(target_arch = "x86")]  
    use core::arch::x86::_rdtsc;  
    #[cfg(target_arch = "x86_64")]  
    use core::arch::x86_64::_rdtsc;  
  
    unsafe { _rdtsc() }  
}
```




How to Improve: Timing

- TimeStamp Counter
 - `constant_tsc`



Without `constant_tsc`




With `constant_tsc`




How to Improve: Timing

- TimeStamp Counter
 - constant_tsc
 - nonstop_tsc



Without nonstop_tsc





With nonstop_tsc



How to Improve: Timing


- TimeStamp Counter
 - `constant_tsc`
 - `nonstop_tsc`
 - Unsynchronized among cores





How to Improve: Timing

- TimeStamp Counter
 - `constant_tsc`
 - `nonstop_tsc`
 - Unsynchronized among cores
 - Threads scheduling




$t_1 < t_2$
 $tsc_1 > tsc_2$





How to Improve: Timing


- TimeStamp Counter
 - `constant_tsc`
 - `nonstop_tsc`
 - Unsynchronized among cores
 - Threads scheduling





How to Improve: Timing


- Synchronize TSCs





How to Improve: Timing


- Synchronize TSCs
 - Offset & Rate
 - `libc::sched_setaffinity()`
 - Retrieve tsc & systime twice





How to Improve: Timing


- Synchronize TSCs
 - Offset & Rate
 - `libc::sched_setaffinity()`
 - Retrieve tsc & systime twice
 - TSC + CPU ID
 - RDTSC + CPUID





How to Improve: Timing

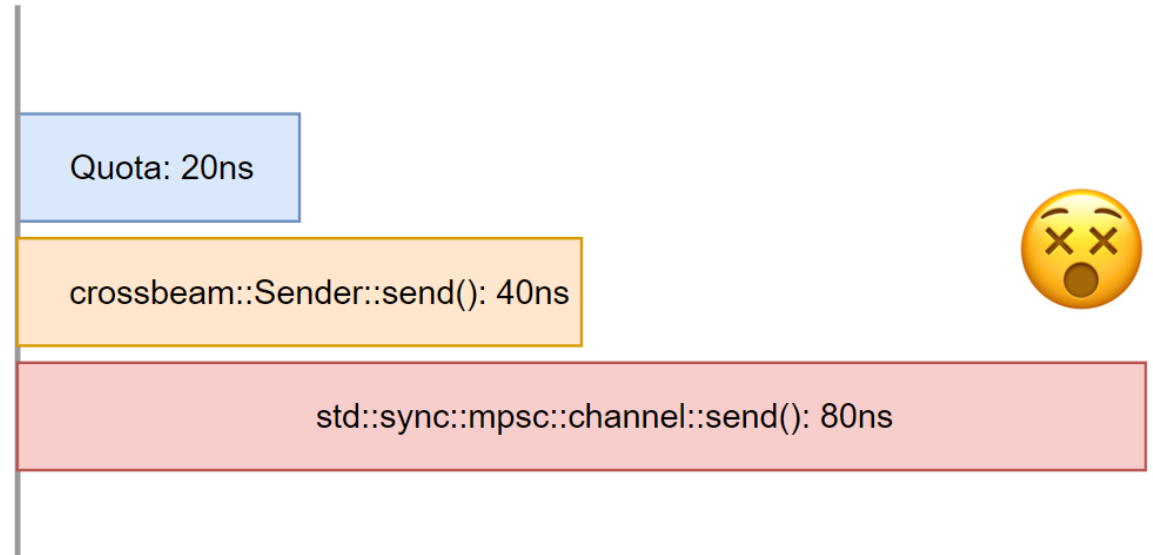
- Synchronize TSCs
 - Offset & Rate
 - `libc::sched_setaffinity()`
 - Retrieve tsc & systime twice
 - TSC + CPU ID
 - RDTSC + CPUID
 - **RDTSCP**





How to Improve: Collection


- Overhead of Span Collection
 - Crossbeam channel
 - Based on atomic variables
 - Prevent compiler from optimizing
 - Unfriendly to CPU cache





How to Improve: Collection

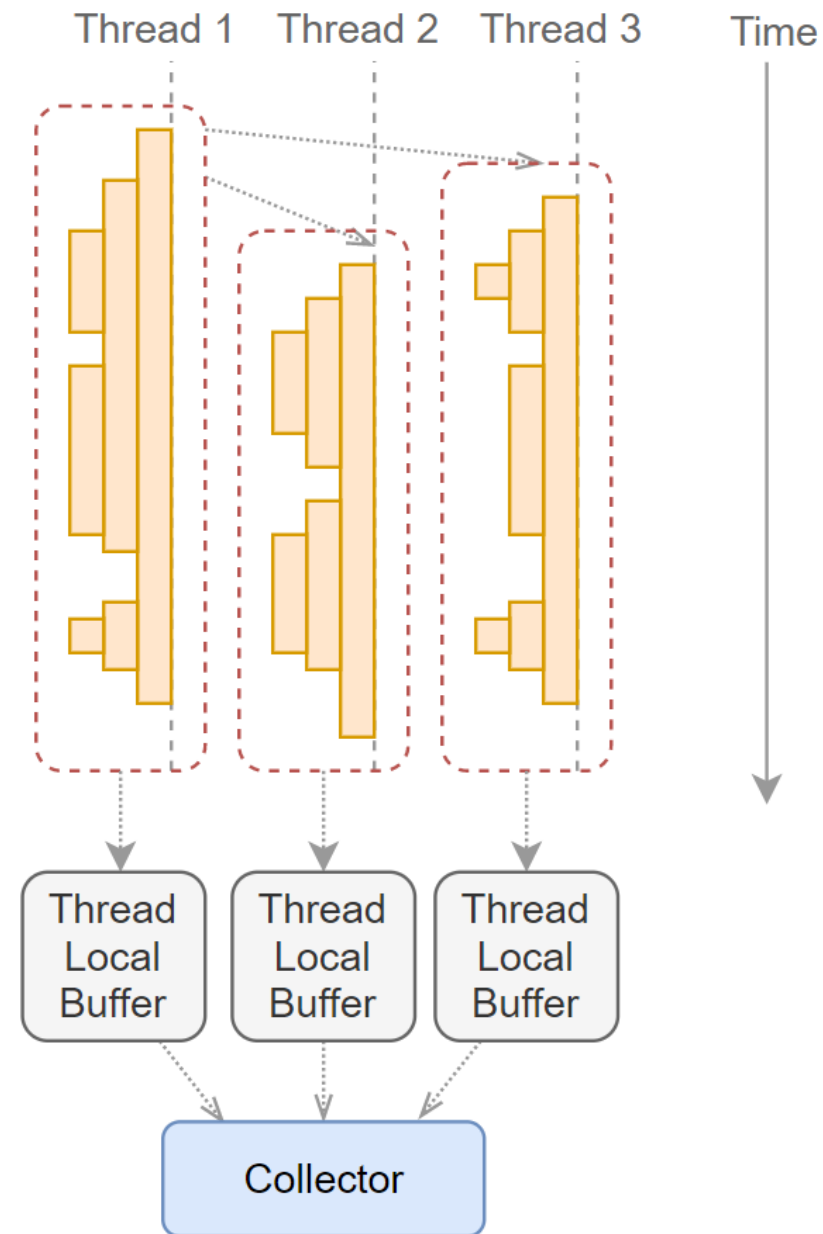
- Improvement: Local + Batch
 - Execution doesn't switch threads all the time





How to Improve: Collection

- Improvement: Local + Batch
 - Execution doesn't switch threads all the time
 - Use **thread local buffer**, collect in batch






How to Improve: Collection

- Improvement: Local + Batch
 - Execution doesn't switch threads all the time
 - Use **thread local buffer**, collect in batch

```
thread_local! {  
    pub static BUFFER: RefCell<Vec<Span>> = RefCell::new(vec![]);  
}
```





Minitrace Usage

- Implicit Context

```
fn foo(ctx: &mut Context) {  
    let span = ctx.create_span();  
    bar(ctx);  
}
```

```
fn bar(ctx: &mut Context) {  
    let span = ctx.create_span();  
}
```

Foo (Parent)

Bar (Child)



Minitrace Usage

- Implicit Context

```
use minitrace::*;  
  
fn foo() {  
+   let _guard = start_span("Foo");  
    bar();  
}  
  
fn bar() {  
+   let _guard = start_span("Bar");  
}
```

Foo (Parent)

Bar (Child)



Minitrace Usage

- Use Macros

```
use minitrace::*;  
use minitrace_macro::*;
```

```
+ #[trace("Foo")]  
fn foo() {  
    bar();  
}
```

```
+ #[trace("Bar")]  
fn bar() { }
```



Minitrace Usage

- Async Function

```
use minitrace::*;  
use minitrace_macro::*;
```

```
+ #[trace_async("Async Foo")]  
async fn async_foo() {  
    async_bar().await;  
}
```

```
+ #[trace_async("Async Bar")]  
async fn async_bar() { }
```



Except Performance

- Safety
 - Unsafe-free (except timing)
 - Thread-safe
 - Thread-local type: !Send, !Sync
- Compatibility with OpenTracing
 - Report to Jaeger and Datadog



Thanks

<https://github.com/tikv/minitrace-rust.git>

<https://github.com/zhongzc>

zhongzhenchi@pingcap.com