



Rust China Conf 2020

Shenzhen, China

2020conf.rustcc.cn



Rust 系统编程在 StratoVirt 中的实践

杨杰

yangjieyj.yang@huawei.com



About Me



- 华为高级工程师
- **StratoVirt** 主要贡献者
- 具备多年 Linux 系统及虚拟化开发经验



What is system programming?


- 任何不是应用程序的软件
 - BIOS、固件、引导加载程序、操作系统内核、hypervisor
- 与硬件交互，直接操作寄存器或内存地址
- 追求高性能和高可靠性
- Assembly -> C



Problems of C

- 无内存管理机制
 - Use-after-free
 - 空指针访问
 - Double-free
 - 内存泄漏
 - 缓冲区溢出
- 语法限制、变量约束不严格
- 不检查数组下标
- 并发访问控制
- 只提供最基础的数据类型和接口

虚拟化组件	开源代码量
libvirt	560 KLOC
QEMU	1570 KLOC
KVM	30 KLOC





Why Rust?

- 内存安全
- 零运行时的性能
- 友好的工具链
- 充满活力的社区 & 生态
-

Use in Industry

- Microsoft
- Facebook
- Amazon
- Google
- DropBox
- Intel and ARM
- CloudFlare
- Mozilla
- Discord

*** Re: Linux kernel in-tree Rust support**
2020-07-10 22:59 ` Josh Triplett
@ 2020-07-10 23:54 ` Linus Torvalds
2020-07-11 21:03 ` Josh Triplett
[0 siblings, 1 reply; 26+ messages in thread](#)
From: Linus Torvalds @ 2020-07-10 23:54 UTC ([permalink](#) / [raw](#))
To: Josh Triplett
Cc: Christian Brauner, Nick Desaulniers, alex.gaynor, Greg KH, geofft, Jbaublitz, Masahiro Yamada, Miguel Ojeda, Steven Rostedt, LKML, clang-built-linux, Kees Cook

On Fri, Jul 10, 2020 at 3:59 PM Josh Triplett <josh@joshtriplett.org> wrote:
>
> As I recall, Greg's biggest condition for initial introduction of this
> was to do the same kind of "turn this Kconfig option on and turn an
> option under it off" trick that LTO uses, so that neither "make
> allnoconfig" nor "make allyesconfig" would require Rust until we've had
> plenty of time to experiment with it.

No, please make it a "is rust available" automatic config option. The exact same way we already do the compiler versions and check for various availability of compiler flags at config time.

See init/Kconfig for things like

```
config LD_IS_LLD
def_bool $(success,$(LD) -v | head -n 1 | grep -q LLD)
```

and the rust support should be similar. Something like

```
config RUST_IS_AVAILABLE
def_bool $(success,$(RUST) ..sometest..)
```

because I _don't_ want us to be in the situation where any new rust support isn't even build-tested by default.

Quite the reverse. I'd want the first rust driver (or whatever) to be introduced in such a simple format that failures will be obvious and simple.

The _worst_ situation to be in is that s (small) group of people start testing their very special situation, and do bad and crazy things because "nobody else cares, it's hidden".



What is StratoVirt?



- 基于 Rust 开发的面向云数据中心的企业级虚拟化平台
- 轻量、安全，实现一套架构统一支持虚拟机、容器、Serverless 三种场景
- 支持 x86_64 和 aarch64

<ul style="list-style-type: none">• StratoVirt:<ul style="list-style-type: none">• 如何开始<ul style="list-style-type: none">• 环境准备• 编译软件• 运行软件• 设计• 如何贡献• 许可	<h2>StratoVirt:</h2> <p>StratoVirt是计算产业中面向云数据中心的企业级虚拟化平台，实现了一套架构统一支持虚拟机、容器、Serverless三种场景。StratoVirt在轻量低噪、软硬协同、Rust语言级安全等方面具备关键技术竞争优势。</p> <p>StratoVirt预留了接口和设计来支持更多特性，未来甚至向标准虚拟化演进。</p> <h3>如何开始</h3>
---	---

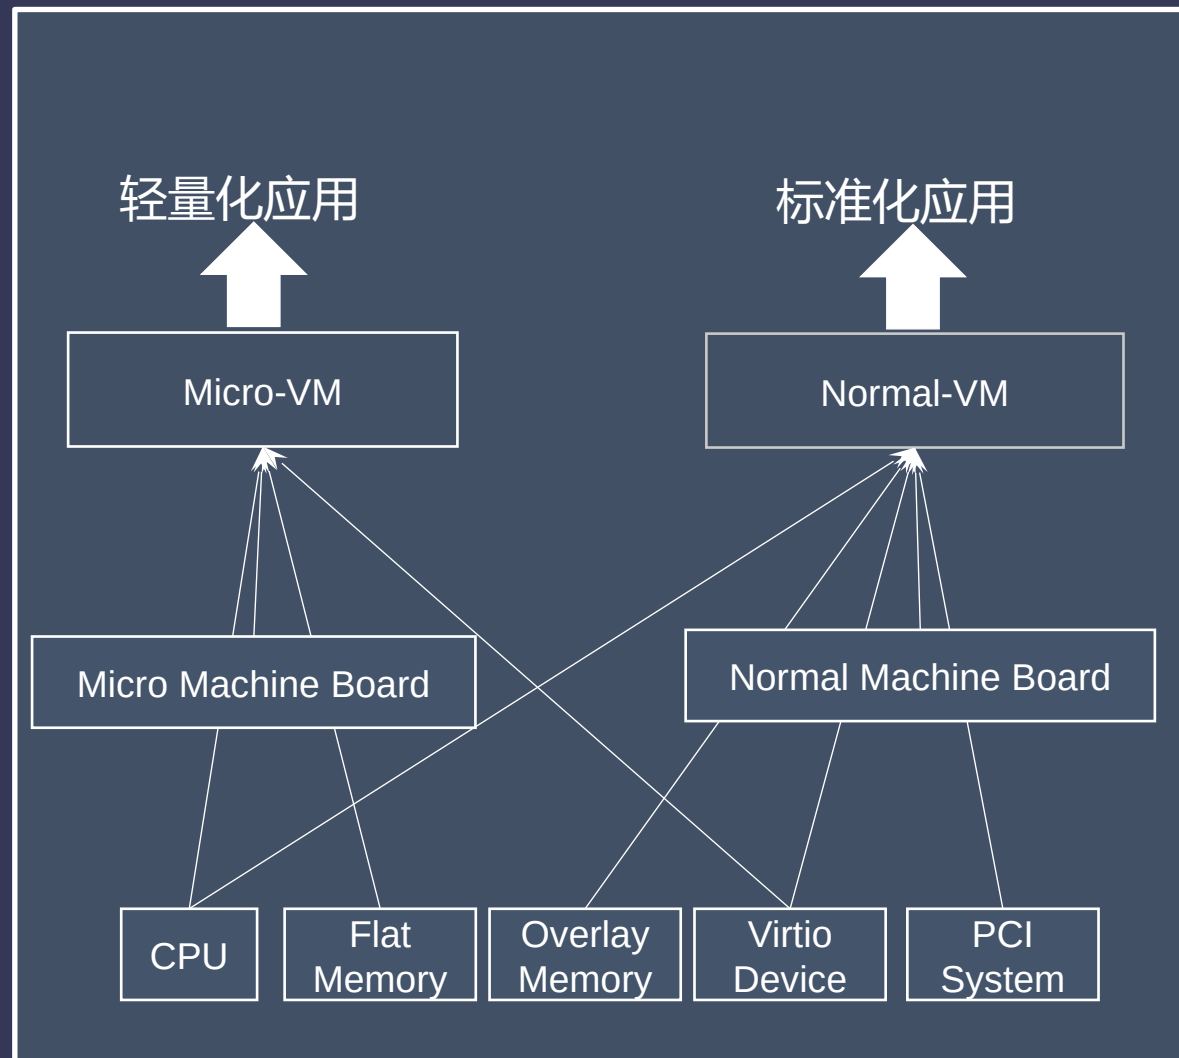
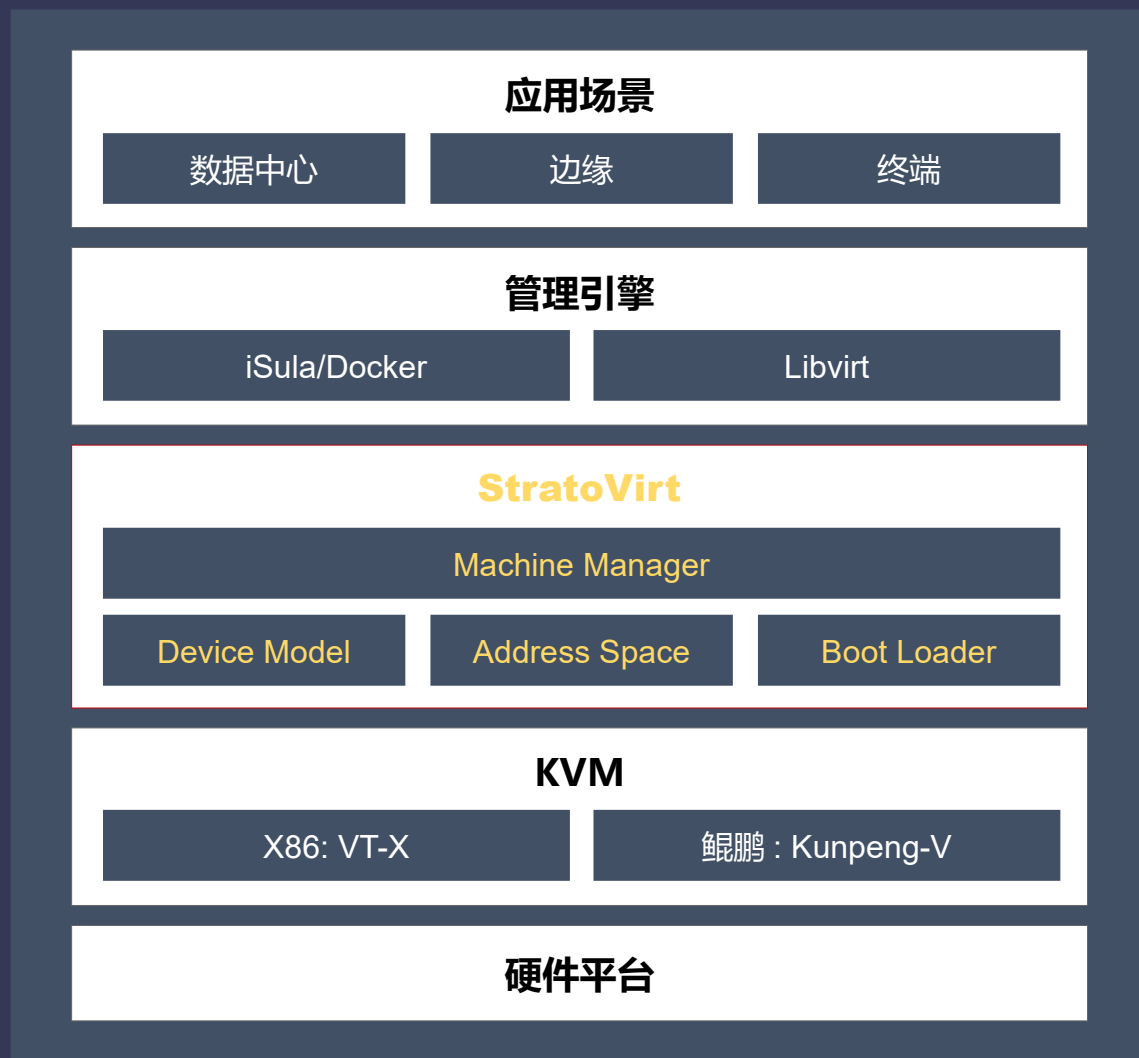


StratoVirt

<https://gitee.com/openeuler/stratovirt>



StratoVirt Architecture





StratoVirt Realization

- 模块化: 1 个 bin + 5 个 lib crate
- 引用第三方 Rust 库 11 个, 涵盖日志、错误处理等
- 支持 libc 的静态链接 (musl) 及动态链接
- 动态链接第三方 C 库
- 30 余个 Linux 系统调用

```
▼ address_space/  
  ▶ src/  
    Cargo.toml  
▼ boot_loader/  
  ▶ src/  
    Cargo.toml  
▼ device_model/  
  ▶ src/  
    Cargo.toml  
▶ docs/  
▶ license/  
▼ machine_manager/  
  ▶ src/  
    Cargo.toml  
▼ src/  
  main.rs  
▶ target/  
▼ util/  
  ▶ src/  
    Cargo.toml  
Cargo.lock  
Cargo.toml  
README.ch.md  
README.md
```



StratoVirt – Virtio-mmio

```
83 /// MMIO Bus.
84 pub struct Bus {
85     /// The devices inserted in bus.
86     devices: Vec<MmioDevice>,
87     /// All replaceable device information.
88     replaceable_info: MmioReplaceableInfo,
89 }
```

```
84 /// MmioDevice structure which used to register into system address space.
85 #[derive(Clone)]
86 pub struct MmioDevice {
87     /// MmioDeviceOps used to be invoked in function realize().
88     device: Arc<Mutex<dyn MmioDeviceOps>>,
89     /// RegionOps used to be registered into system address space.
90     region_ops: RegionOps,
91     /// The DeviceResource required by this MMIO device.
92     resource: Arc<DeviceResource>,
93 }
```

```
188 /// Trait for MMIO device.
189 pub trait MmioDeviceOps: Send + DeviceOps {
190     /// Realize this MMIO device for VM.
191     fn realize(&mut self, vm_fd: &VmFd, resource: DeviceResource) -> Result<()>;
192
193     /// Get the resource requirement of MMIO device.
194     fn get_type(&self) -> DeviceType;
195
196     /// Update the low level config of MMIO device.
197     fn update_config(&mut self, _dev_config: Option<Arc<dyn ConfigCheck>>) -> Result<()> {
198         bail!("Unsupported to update configuration");
199     }
200
201     /// Get IoEventFds of MMIO device.
202     fn ioeventfds(&self) -> Vec<RegionIoEventFd> {
203         Vec::new()
204     }
205 }
```

vm

Virtio-mmio



Virtio-mmio

Virtio-blk

Virtio-net

Virtio-rng

StratoVirt

```
292 /// virtio-mmio device structure.
293 pub struct VirtioMmioDevice {
294     /// The entity of low level device.
295     device: Arc<Mutex<dyn VirtioDevice>>,
296     /// Identify if this device is activated by frontend driver.
297     device_activated: bool,
298     /// EventFd used to send interrupt to VM
299     interrupt_evt: EventFd,
300     /// HostNotifyInfo used for guest notifier
301     host_notify_info: HostNotifyInfo,
302     /// Virtio common config refer to Virtio Spec.
303     common_config: VirtioMmioCommonConfig,
304     /// System address space.
305     mem_space: Arc<AddressSpace>,
306 }
```



StratoVirt – IO handling

- 单线程 epoll：满足基本的 IO 性能需求，降低内存开销
- 存储设备后端支持异步 IO

```
96 /// MainLoop manager, advise continue running or stop running
97 pub trait MainLoopManager {
98     fn main_loop_should_exit(&self) -> bool;
99     fn main_loop_cleanup(&self) -> Result<()>;
100 }
101
102 /// Main Epoll Loop Context
103 #[allow(clippy::vec_box)]
104 pub struct MainLoopContext {
105     /// Epoll file descriptor.
106     epoll: Epoll,
107     /// Control epoll loop running.
108     manager: Option<Arc<dyn MainLoopManager>>,
109     /// Fds registered to the 'MainLoop'.
110     events: Arc<RwLock<BTreeMap<i32, Box<EventNotifier>>>>,
111     /// Events abandoned are stored in garbage collector.
112     gc: Arc<RwLock<Vec<Box<EventNotifier>>>>,
113     /// Temp events vector, store wait returned events.
114     ready_events: Vec<EpollEvent>,
115 }
116
117
118 pub type NotifierCallback = dyn Fn(EventSet, RawFd) -> Option<Vec<EventNotifier>>;
119
120 /// Epoll Event Notifier Entry.
121 pub struct EventNotifier {
122     /// Raw file descriptor
123     pub raw_fd: i32,
124     /// Notifier operation
125     pub op: NotifierOperation,
126     /// Parked fd, temporarily removed from epoll
127     pub parked_fd: Option<i32>,
128     /// The types of events for which we use this fd
129     pub event: EventSet,
130     /// Event Handler List, one fd event may have many handlers
131     pub handlers: Vec<Arc<Mutex<Box<NotifierCallback>>>>,
132     /// Event status
133     status: EventStatus,
134 }
```

```
82 pub fn handle(&mut self) -> Result<()> {
83     let evts = self.ctx.get_events()?;
84     for e in evts.events.iter().take(evts.nr) {
85         if e.res2 == 0 {
86             unsafe {
87                 let node = e.data as *mut CbNode<T>;
88
89                 (self.complete_func)((*node).value, e.res);
90                 self.aio_in_flight.unlink(&(*node));
91
92                 // free mem
93                 if let Some(i) = (*node).value.iocb {
94                     libc::free((*node).value.iovec.as_ptr() as *mut libc::c_void);
95                     libc::free(i.as_ptr() as *mut libc::c_void);
96                 };
97                 libc::free(node as *mut libc::c_void);
98             }
99         }
100     }
101     self.process_list()
102 }
103
104 fn process_list(&mut self) -> Result<()> {
105     if self.aio_in_queue.len > 0 && self.aio_in_flight.len < self.max_events {
106         let mut iocbs = Vec::new();
107
108         for _ in self.aio_in_flight.len..self.max_events {
109             match self.aio_in_queue.pop_tail() {
110                 Some(node) => {
111                     iocbs.push(node.value.iocb.unwrap().as_ptr());
112                     self.aio_in_flight.add_head(node);
113                 }
114                 None => break,
115             }
116         }
117
118         if !iocbs.is_empty() {
119             return self.ctx.submit(iocbs.len() as i64, &mut iocbs);
120         }
121     }
122
123     Ok(())
124 }
```



StratoVirt – Error handling

- 提高可维性
- 代码优雅、简洁

```
73 pub mod errors {
74     error_chain! {
75         links {
76             AddressSpace(address_space::errors::Error, address_space::errors::ErrorKind);
77             Util(util::errors::Error, util::errors::ErrorKind);
78             BootLoader(boot_loader::errors::Error, boot_loader::errors::ErrorKind);
79             Manager(machine_manager::errors::Error, machine_manager::errors::ErrorKind);
80             Cpu(crate::cpu::errors::Error, crate::cpu::errors::ErrorKind);
81             Mmio(crate::mmio::errors::Error, crate::mmio::errors::ErrorKind);
82         }
83         foreign_links {
84             Io(std::io::Error);
85             Kvm(kvm_ioctls::Error);
86             Json(serde_json::Error);
87             Nul(std::ffi::NulError);
88         }
89     }
90 }
```

```
42 pub mod errors {
43     error_chain! {
44         links {
45             AddressSpace(address_space::errors::Error, address_space::errors::ErrorKind);
46             Virtio(crate::virtio::errors::Error, crate::virtio::errors::ErrorKind);
47         }
48         errors {
49             MmioRegister(offset: u64) {
50                 display("Unsupported mmio register, 0x{:x}", offset)
51             }
52             DeviceStatus(status: u32) {
53                 display("Invalid device status 0x{:x}", status)
54             }
55         }
56     }
57 }
```



Challenges

- Unsafe
 - 调用 C 接口
 - 性能敏感
 - 反序列化
 - 访问全局可变变量
 - 链表



Challenges

- Trait object

E0225

Run

Multiple types were used as bounds for a closure or trait object.

Erroneous code example:

```
i fn main() {  
    let _: Box<dyn std::io::Read + std::io::Write>;  
}
```

Rust does not currently support this.

Auto traits such as `Send` and `Sync` are an exception to this rule: It's possible to have bounds of one non-builtin trait, plus any number of auto traits. For example, the following compiles correctly:


```
fn main() {  
    let _: Box<dyn std::io::Read + Send + Sync>;  
}
```

```
83 83  /// MmioDevice structure which used to register into system address space.  
84 84  #[derive(Clone)]  
85 85  pub struct MmioDevice {  
86 86      /// MmioDeviceOps used to be invoked in function realize().  
87 87      device: Arc<Mutex<dyn MmioDeviceOps>>,  
88 88      /// RegionOps used to be registered into system address space.  
89 88      - dev_region: Arc<Mutex<dyn RegionOps>>,  
89 89      + region_ops: RegionOps,  
90 90      /// The DeviceResource required by this MMIO device.  
91 91      resource: Arc<DeviceResource>,  
92 92  }  
93 93
```



Challenges

- 锁
 - 性能
 - 努力避免死锁





Challenges

- 第三方 crate
 - 可靠
 - 稳定
 - 活跃
 - 满足功能需求



Challenges

- 热补丁技术

```
x-1 /mnt/codes/qvisor/qvisor/qvisor [master|✓]
10:18 $ nm target/debug/stratovirt | grep realize
000000000012e360 T _ZN3cpu6x86_646X86CPU7realize17h58cabe905cfd6d1eE
000000000013aab0 T _ZN46_$LT$cpu..CPU$u20$as$u20$cpu..CPUInterface$GT$7realize17hf6dea4a6a61c701cE
00000000000c3370 T _ZN4mmio10MmioDevice7realize17h058dce6e21f80939E
00000000000b9700 T _ZN4mmio3bus3Bus15realize_devices17h42d8514e36842b60E
000000000026ded0 T _ZN4util7seccomp13SyscallFilter7realize17h86c3a1b5acbaafdbE
00000000000f4c40 T _ZN57_$LT$virtio..net..Net$u20$as$u20$virtio..VirtioDevice$GT$7realize17h490a5531f09ce566E
0000000000011f0c0 t _ZN57_$LT$virtio..net..Net$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h11e95b44ed788801E
0000000000011f0b0 t _ZN57_$LT$virtio..net..Net$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17hcb589451c99d8701E
00000000000110720 T _ZN61_$LT$virtio..block..Block$u20$as$u20$virtio..VirtioDevice$GT$7realize17h5102522ebdb3603E
00000000000116270 t _ZN61_$LT$virtio..block..Block$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h05e7940caaeb9fbceE
000000000001161b0 t _ZN61_$LT$virtio..block..Block$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h143f16b7cacee153E
000000000001160f0 t _ZN61_$LT$virtio..block..Block$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h261181caa76e880feE
000000000001160e0 t _ZN61_$LT$virtio..block..Block$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17hcd78509fb52310c0E
00000000000b16c0 T _ZN62_$LT$legacy..serial..Serial$u20$as$u20$mmio..MmioDeviceOps$GT$7realize17h44a658a54e22cec6E
00000000000b2240 t _ZN62_$LT$legacy..serial..Serial$u20$as$u20$mmio..MmioDeviceOps$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h65537cc0f3a34e7dE
00000000000e0f20 T _ZN65_$LT$virtio..console..Console$u20$as$u20$virtio..VirtioDevice$GT$7realize17h5792272600c6d999E
00000000000118890 T _ZN72_$LT$virtio..vhost..kernel..net..Net$u20$as$u20$virtio..VirtioDevice$GT$7realize17haa775ca30c7335faE
00000000000124160 t _ZN72_$LT$virtio..vhost..kernel..net..Net$u20$as$u20$virtio..VirtioDevice$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17hc7ba587b21aa6312E
00000000000c7450 T _ZN75_$LT$mmio..virtio_mmio..VirtioMmioDevice$u20$as$u20$mmio..MmioDeviceOps$GT$7realize17h437481f702fe011aE
00000000000c2740 t _ZN75_$LT$mmio..virtio_mmio..VirtioMmioDevice$u20$as$u20$mmio..MmioDeviceOps$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h2b4d0ada48b6c269E
00000000000c2750 t _ZN75_$LT$mmio..virtio_mmio..VirtioMmioDevice$u20$as$u20$mmio..MmioDeviceOps$GT$7realize28_$u7b$$u7b$closure$u7d$$u7d$17h5fd27502041511e3E
00000000000eefe0 T _ZN76_$LT$virtio..vhost..kernel..vsock..Vsock$u20$as$u20$virtio..VirtioDevice$GT$7realize17hb0a7b6842d67fcdceE
000000000008c0f0 T _ZN8micro_vm7machine12LightMachine7realize17h4c956ff9a17ba601E
```



**Rust is still *young*, but
*promising***